



Tips for Linux Kernel Development

Jessica Yu, Red Hat
UC Santa Cruz CMPS107
March 7, 2017

Overview

- How do I get started?
- Personal Experiences
- Tips for speeding up compilation
- Debugging tutorial with QEMU

How do I get started?

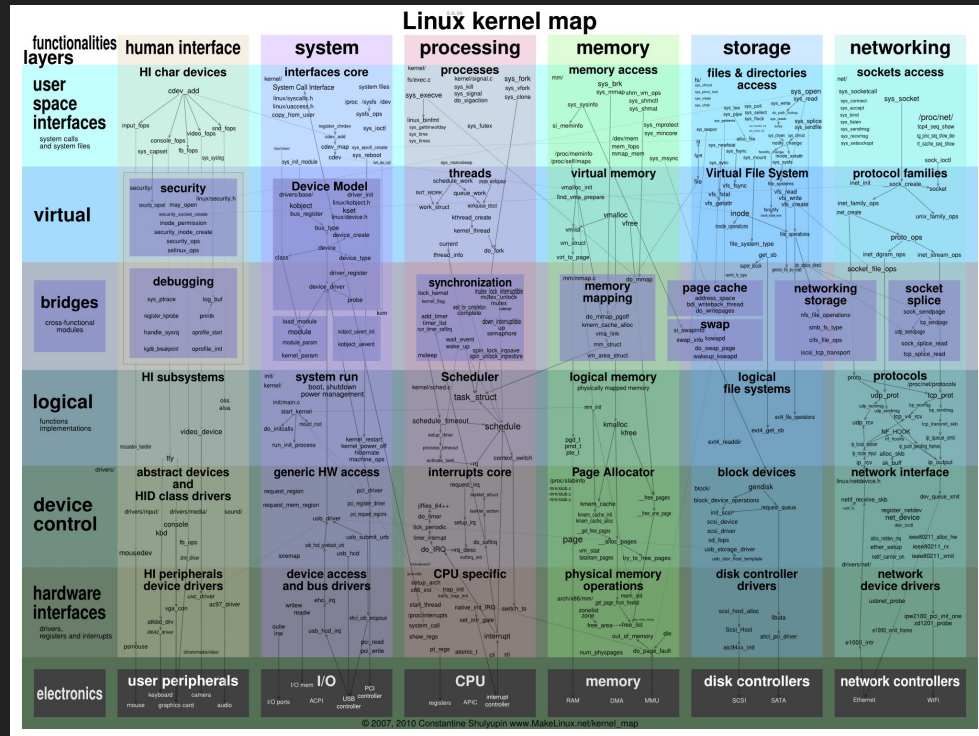




Image source: http://www.makelinux.net/kernel_map/LKM3_2048.png

How do I get started?

- Compile your own kernel and boot it! (in a VM, on your laptop, etc)
 - If you boot it on your own machine, make sure to have backup kernels :-)
- Stick a print statement somewhere in the kernel source
 - (in a function you know will run), compile, boot, and look for it in dmesg
- Write a trivial kernel module and load it (~15 lines)
- The Eudryptula Challenge
- Linux kernel internships for university students
 - Google Summer of Code
 - TONS of open source projects! Mozilla, FreeBSD, Linux Foundation, Coreboot...
 - Outreachy (formerly called Outreach Program for Women)
 - Similar to GSoC, contribute to open source projects (incl. Linux kernel!)
 - Also open to non-students
 - **Both GSoC and Outreachy are still accepting applications!!**

How do I get started?

- Find a bug and fix it
 - <http://kernelnewbies.org/FoundBug>
 - <https://bugzilla.kernel.org/>
- Lurk on various mailing lists
 - Get a sense of how patches are accepted, how to respond to patch comments, etc.
 - All the vger mailing lists: <http://vger.kernel.org/vger-lists.html>
 - E.g. linux-ext4, live-patching, linux-doc
- LWN.net is a great resource!
- Books
 - Linux Kernel Development (Robert Love)
 - Understanding the Linux Kernel (Daniel Bovet, Marco Cesati)



The Eudryptula Challenge

- A series of tasks related to the Linux kernel, geared towards familiarizing you with Linux kernel development
 - starts out easy...gets more complex
 - write a simple kernel module
 - write a syscall
 - send in a patch fixing coding style
 - etc...
 - 20 tasks in total
 - I hear there's a secret prize at the end!

First attempt

- Task 10 of the Eudryptula Challenge.....send a real patch!
 - "Create a patch that fixes one coding style problem in any of the files in drivers/staging/"

Jessica Yu | 28 Jul 05:53 2014

[PATCH] Staging: lustre: linux-module: fix coding style issues.

Fixed some coding style issues.

Signed-off-by: Jessica Yu <jyu <at> cowsay.org>

drivers/staging/lustre/lustre/libcfs/linux/linux-module.c | 14 ++++++-----

1 file changed, 7 insertions(+), 7 deletions(-)

First attempt

Greg Kroah-Hartman | 27 Jul 22:01 2014



Re: [PATCH] Staging: lustre: linux-module: fix coding style issues.

On Sun, Jul 27, 2014 at 08:53:50PM -0700, Jessica Yu wrote:

> Fixed some coding style issues.

What coding style issues? Be specific, and explicit.

Care to try it again?

greg k-h

Take two



gregkh@linuxfoundation.org

30.07.14



an jyu 

This is a note to let you know that I've just added the patch titled

Staging: lustre: linux-module: remove unnecessary spaces

to my staging git tree which can be found at

git://git.kernel.org/pub/scm/linux/kernel/git/gregkh/staging.git
in the staging-next branch.

Hurray!!!

Tips for speeding up compilation



Tips for speeding up compilation

- **ccache**
 - cache compiled objects
 - greatly speeds up recompilation
- **Parallel compilation**
 - `make -j 16` on a system with 8 cores
 - Good rule of thumb: `make -j $(number of cores)`
- **In-memory compilation (tmpfs)**
 - Avoids filesystem syncing/flushing
 - Recommended only if you have enough RAM to spare...
 - Beware of OOM situations
 - ``free -h`` tells you total/used/free memory
- **All this combined**
 - `make -j 16 O=/tmp/linux 271,09s user 195,99s system 550% cpu 1:24,85 total`
 - (Intel Core i7-4800MQ CPU @ 2.70GHz and 16GB of DDR3 RAM)

Tips for speeding up compilation

- Minimize kernel config
 - A lot of time is spent compiling kernel modules!
 - Avoid compiling modules you know you don't need
 - Kernels provided by distributions (i.e. Ubuntu, etc.) tend to be more generalized
 - (provide the config that will work on most machines)
 - means a lot of modules are included
 - ``make localmodconfig`` can help slim down your `.config`
 - uses `lsmod` to create a config with only the currently loaded modules enabled

Tips for debugging

- Demo time!
 - QEMU + gdb

Debugging with QEMU

- What is QEMU?
 - A virtualization tool like VirtualBox
 - allows you to run a complete operating system on your machine.
 - very useful for trying out different OS's, testing software, and running applications that won't run on your machine's native platform.
- Get QEMU: <http://wiki.qemu.org/Download>
 - Your distribution more than likely already provides a package for qemu
- We can use QEMU and gdb to help us debug kernels

Debugging with QEMU

- Compiling a kernel for QEMU
 - `make olddefconfig`
 - ``make help`` to see all config targets
 - enable `CONFIG_DEBUG_INFO`
 - `make -j ...`

Debugging with QEMU

- Prepare your kernel for gdb debugging...
 - Turn on debug info in your .config => CONFIG_DEBUG_INFO

```
.config - Linux/x86 4.5.0-rc6 Kernel Configuration
-> Kernel hacking -> Compile-time checks and compiler options
Compile the kernel with debug info

CONFIG_DEBUG_INFO:

If you say Y here the resulting kernel image will include
debugging info resulting in a larger kernel image.
This adds debug symbols to the kernel and modules (gcc -g), and
is needed if you intend to use kernel crashdump or binary object
tools like crash, kgdb, LKCD, gdb, etc on the kernel.
Say Y here only if you plan to debug the kernel.

If unsure, say N.

Symbol: DEBUG_INFO [=n]
Type : boolean
Prompt: Compile the kernel with debug info
Location:
-> Kernel hacking
-> Compile-time checks and compiler options
Defined at lib/Kconfig.debug:120
Depends on: DEBUG_KERNEL [=y] && !COMPILE_TEST [=n]
```


Debugging with QEMU

```
$ qemu-system-x86_64 -s -S \ (1)
    -kernel /path/to/vmlinuz \ (2)
    -initrd initramfs.cpio.gz \ (3)
    -nographic \ (4)
    -append "console=ttyS0" (5)
```

1. **-s** Shorthand for `-gdb tcp::1234`, i.e. QEMU listen on TCP port 1234 for a connection by gdb.
-S Do not start CPU at startup. QEMU will start as if you set a breakpoint at time zero, and waits for gdb to connect
2. Path to kernel
3. Path to initramfs (we'll just be booting into a very minimal userland, using BusyBox)

Debugging with QEMU

```
$ qemu-system-x86_64 -s -S \ (1)
  -kernel /path/to/vmlinuz \ (2)
  -initrd initramfs.cpio.gz \ (3)
  -nographic \ (4)
  -append "console=ttyS0" (5)
```

4. **-nographic -append "console=ttyS0"**

- a. Disable graphical output so that QEMU is a simple command line application.
- b. All kernel output will go to the ttyS0 serial console, with will be printed to stdio (your terminal).

Debugging with QEMU

- Attaching gdb to QEMU
 - `$ cgdb -ex "target remote localhost:1234" vmlinux`
 - cgdb is just the curses interface to gdb. It's prettier. :-)
 - **-ex** Execute specified gdb command.
 - **target remote localhost:1234** Establish a tcp connection to port 1234 on host 'localhost' (where the QEMU gdb stub is listening!)

Debugging with QEMU

- Kernel debugging demo
 - set breakpoints at two functions
 - `cmdline_proc_show()`
 - `fs/proc/cmdline.c`
 - This function is called when we read from `/proc/cmdline`
 - `meminfo_proc_show()`
 - `fs/proc/meminfo.c`
 - This function is called when we read from `/proc/meminfo`

Debugging with QEMU

- Debugging NULL dereference demo

```
[    7.186020] BUG: unable to handle kernel NULL pointer dereference at (null)
[    7.186706] IP: cmdline_proc_show+0x17/0x30  <--
...
[    7.187005] Oops: 0002 [#1] SMP
[    7.187005] Modules linked in:
[    7.187005] CPU: 0 PID: 111 Comm: cat Not tainted 4.10.1+ #10
[    7.187005] Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS
1.8.2-20150714_191134- 04/01/2014
[    7.187005] task: fffff8800065d0000 task.stack: fffffc90000254000
[    7.187005] RIP: 0010:cmdline_proc_show+0x17/0x30
```

Then in gdb...

```
(gdb) list *(cmdline_proc_show+0x17)
0xffffffff812de257 is in cmdline_proc_show (fs/proc/cmdline.c:10).
...
```

Debugging with QEMU

- We were able to debug and step through kernel code
- What about linux kernel modules?
 - trickier, but doable
- Modules are **dynamically** loaded
 - gdb won't know where the module will be loaded
 - tell gdb about the module after module load
 - `add-symbol-file` gdb command
 - requires section addresses for:
 - `.text`
 - `.data`
 - `.bss`
 - A module's section addresses can be found under `/sys/module/<module_name>/sections`

Debugging with QEMU

- gdb won't know about your module's symbols unless you gdb about it
- use gdb command **add-symbol-file** to give gdb additional symbol information for debugging

```
(gdb) help add-symbol-file
```

```
Load symbols from FILE, assuming FILE has been dynamically loaded.
```

```
Usage: add-symbol-file FILE ADDR [-s <SECT> <SECT_ADDR> ...]
```

```
ADDR is the starting address of the file's text.
```

```
(gdb) add-symbol-file /path/to/module.ko <text_addr> -s .data <data_addr> -s  
.bss <bss_addr>
```

- Section addresses are found under `/sys/kernel/<module>/sections`

```
$ cat /sys/module/cmeps107/sections/.text
```

```
0xfffffffffa000000
```

```
$ cat /sys/module/cmeps107/sections/.data
```

```
0xfffffffffa0000c0
```

Debugging with QEMU

- Sample module: `cmps107.ko`
 - Demo code: <https://github.com/flaming-toast/cmps107-demo>
- `cmps107` module does the following:
 - Creates subdirectory "`cmps107`" under `/sys/kernel/`
 - Creates file "`foo`" under `/sys/kernel/cmps107/`
- The functions we want to step through in `gdb`:
 - `foo_show()` - called when file `foo` is read from
 - `foo_store()` - called when file `foo` is written to

Debugging with QEMU

- Debugging kernel module demo

Debugging with QEMU

- More useful walkthroughs:
 - <http://files.meetup.com/1590495/debugging-with-qemu.pdf>
 - <http://mgalgs.github.io/2015/05/16/how-to-build-a-custom-linux-kernel-for-qemu-2015-edition.html>
 - <http://wiki.osdev.org/QEMU>

Other tips

- Navigating the kernel source
 - The Linux Cross Reference (LXR) - <http://lxr.free-electrons.com/>
 - Very useful for finding definitions of functions, structs, where a function is called, etc
 - **ctags** with vim
 - jump to function or struct definitions
 - More tips here: <http://kernelnewbies.org/FAQ/CodeBrowsing>
- Submitting your first patch to lkml
 - Greg Kroah-Hartman has an excellent tutorial on this:
 - <https://www.youtube.com/watch?v=LLBrBBImJt4>

Google Summer of Code

- Google Summer of code <https://summerofcode.withgoogle.com/>
 - Applications **open** on March 20, 2017, and **close** on April 3, 2017
 - List of participating organizations this summer:
<https://summerofcode.withgoogle.com/organizations/>
 - **The Linux Foundation's** GSoC page for this year (list of project ideas for students):
<https://wiki.linuxfoundation.org/gsoc/google-summer-code-2017>
 - Participating GSoC organizations related to **operating systems**:
https://summerofcode.withgoogle.com/organizations/?sp-category=operating_systems

Red Hat summer internships

- Program runs from mid-May to mid-August
 - Opportunities to work on the kernel as an intern
 - List of engineering internship positions open:
 - <https://careers-redhat.icims.com/jobs/search?ss=1&searchKeyword=internship&searchCategory=17505>
 - Drop me an email with resume if you're interested!

Thanks!

Got questions? Feel free to drop me an email!

jeyu at redhat.com

More useful links!

- **Navigating the kernel source**
 - The Linux Cross Reference (LXR) - <http://lxr.free-electrons.com/>
 - <http://kernelnewbies.org/FAQ/CodeBrowsing>
- **Submitting your first patch to lkml**
 - Greg Kroah-Hartman's tutorial: <https://www.youtube.com/watch?v=LLBrBBImJt4>
- **QEMU tutorials**
 - <http://mgalgs.github.io/2015/05/16/how-to-build-a-custom-linux-kernel-for-qemu-2015-edition.html>
 - <http://files.meetup.com/1590495/debugging-with-qemu.pdf>
- **Filesystem images you can use with QEMU:** <http://fs.devloop.org.uk/>
- **Debugging your kernel**
 - <https://wiki.ubuntu.com/Kernel/KernelDebuggingTricks>
 - http://wiki.osdev.org/Kernel_Debugging
- **Debugging linux kernel modules**
 - <https://www.linux.com/learn/kernel-newbie-corner-kernel-and-module-debugging-gdb>
- **Participating in the Linux kernel community**
 - <http://www.linuxfoundation.org/content/how-participate-linux-community>
- **Creating a custom initramfs:** https://wiki.gentoo.org/wiki/Custom_Initramfs